

Zero to Zero.Five

From Zero to Zero.Five — Hardening My First Homelab with Proxmox, Tailscale, and SSH Security

Over the past few weeks, I took a dusty old gaming PC and transformed it into a secure, purpose-built homelab. Not just to spin up random VMs — but to create something **reliable, hardened, and production-like**, where I could simulate real-world identity, networking, and red-team scenarios in a safe environment.

This write-up documents that process — from installing Proxmox, to securing remote access with Tailscale, to enforcing zero-trust principles through firewalls, ACLs, and SSH key management.

The goal? Treat this like an actual enterprise asset, even if it's just running under my desk.

Installing Proxmox Virtual Environment

I started with [Proxmox VE](#) — a Type 1 hypervisor built on Debian, widely used for homelabs due to its GUI, container support, and KVM-based virtualization.

The install process was simple:

1. Created a bootable USB with Rufus
2. Installed Proxmox on bare metal
3. Configured a hostname, root password, and static IP



Welcome to Proxmox Virtual Environment

Install Proxmox VE (Graphical)
Install Proxmox VE (Terminal UI)
Install Proxmox VE (Terminal UI, Serial Console)
Advanced Options

enter: select, arrow keys: navigate, e: edit entry, esc: back

From here, I was dropped into a full-featured hypervisor. But before I started deploying anything, I knew I needed to address the elephant in the room:

Default Security Gaps in Proxmox

By default, Proxmox prioritizes usability — not security. That's fine for internal-only labs, but mine was going to support simulated attacks and remote admin access. Some major concerns:

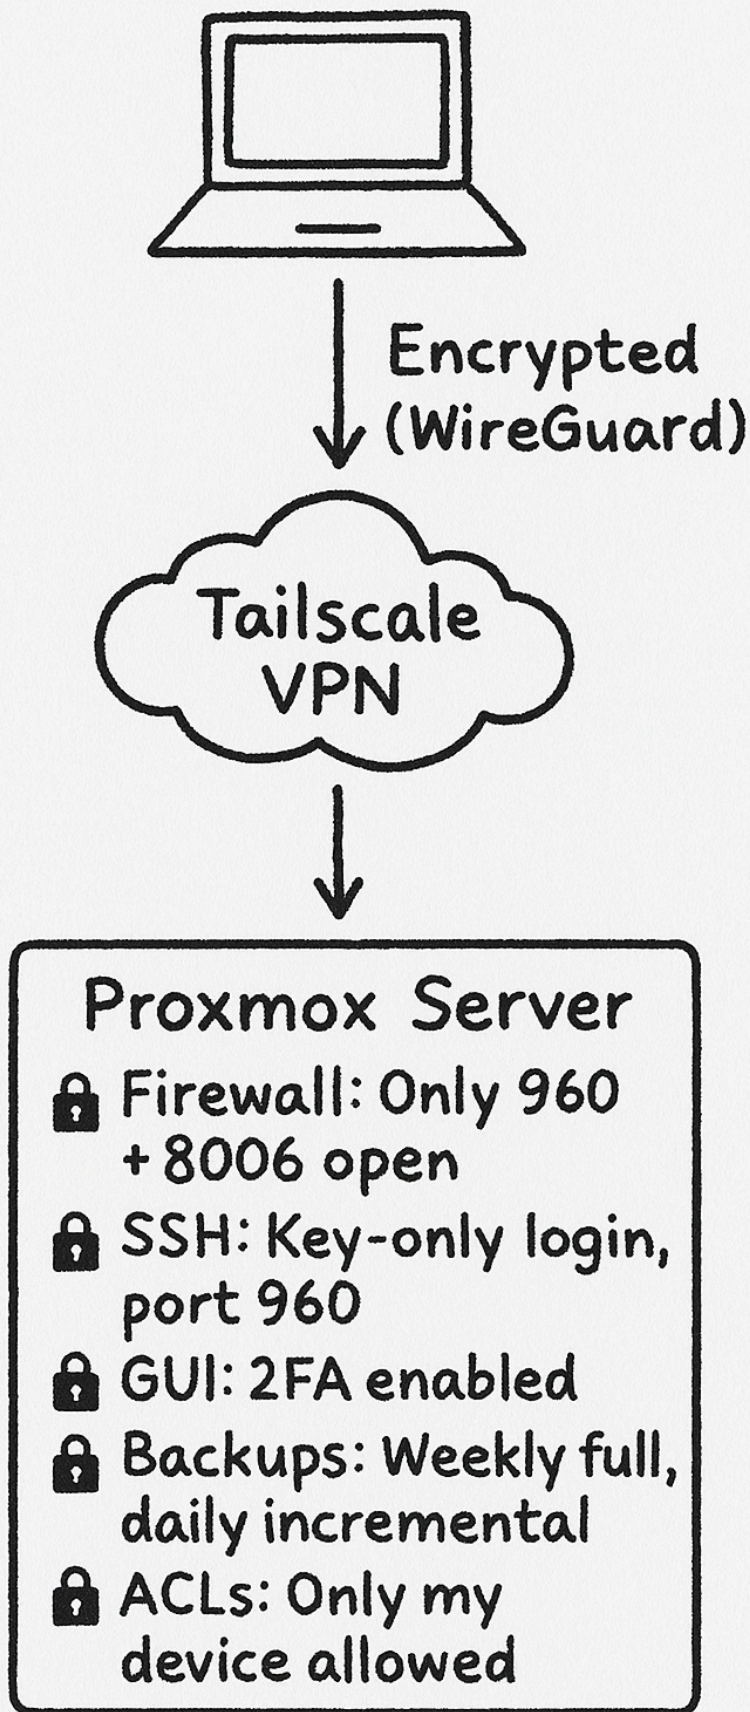
- **Root login via SSH was allowed** — attackers could directly target the most privileged account
- **SSH password authentication enabled** — makes brute-force login attempts trivial
- **No multi-factor authentication (MFA)** for the web UI
- **Proxmox web GUI exposed on default port 8006**, often left open to the internet

These are all commonly exploited misconfigurations in real-world breaches. So instead of rushing to create VMs, I decided to take a “secure-first” approach and lock the environment down immediately.

Private Remote Access with Tailscale

In enterprise networks, VPNs are a standard requirement for accessing anything sensitive. At previous jobs, I'd seen solutions like GlobalProtect or Cisco AnyConnect. I wanted a similar experience — but lightweight and identity-aware.

I initially tried WireGuard, but quickly pivoted to **Tailscale** — a drop-in VPN built on WireGuard that takes care of networking, firewalls, and DNS for you.



Why Tailscale?

- **Peer-to-peer encrypted mesh** — faster and more secure than hub-and-spoke models
- **Identity-aware access** — enforces access based on who you are, not just your IP

- **Zero port forwarding** — no need to expose my router to the world
- **Extremely easy setup** — literally `tailscale up`, log in with Google, and you're online

This took my Proxmox GUI from "**visible to the internet**" to "**only reachable through an encrypted tunnel authenticated to my identity.**"

🔥 Basic Firewall with UFW

Even though Tailscale effectively “hides” the machine from the public internet, I wanted to go a step further. **Defense in depth** is the name of the game. If something fails (like Tailscale disconnecting), I don't want SSH or the GUI suddenly accessible.

So I installed **UFW** (Uncomplicated Firewall) and set strict rules:

bash

CopyEdit

```
sudo ufw allow 960/tcp # SSH (custom port) sudo ufw allow 8006/tcp # Proxmox GUI sudo  
ufw default deny incoming sudo ufw default allow outgoing sudo ufw enable
```

```
apt install ufw  
ufw allow 960/tcp      # Allow SSH on custom port  
ufw allow 8006/tcp     # Allow Proxmox Web GUI  
ufw default deny incoming  
ufw default allow outgoing  
ufw enable  
ufw status
```

Why This Matters:

- Reduces my server's network “attack surface”
- Keeps services reachable **only** from trusted sources (Tailscale peers)
- Protects against accidental exposure if I ever change my networking config

I also changed SSH from the default port 22 to 960 — not true security by itself, but it dramatically reduces automated scan attempts from bots that blindly target 22.

🚫 Killing Root Login and Using Least Privilege

Direct root access over SSH is dangerous. Even in small setups, it's best to avoid logging in as root entirely. Instead, I:

- Created a new user (`denney`)
- Gave it sudo access for administrative tasks
- Disabled SSH login for `root` in `/etc/ssh/sshd_config`

Why?

- Limits the “blast radius” if the account is compromised
- Enforces better operational habits (privilege escalation only when needed)
- Aligns with real-world practices (nobody logs in as root in modern enterprise)

🔑 SSH Key-Based Authentication

Passwords are the weakest link in any authentication chain. Even strong ones can be phished or brute-forced. So I replaced SSH password login entirely with **public-key authentication**.

```
ssh-keygen -t ed25519 -C "your_email@example.com"
```

```
Generating public/private ed25519 key pair.  
Enter file in which to save the key (/c/Users/yourname/.ssh/id_ed25519):  
|
```

Steps taken:

- Generated an ED25519 key pair
- Copied the public key to the server's `~/.ssh/authorized_keys`
- Disabled password login in `sshd_config`: `nginx CopyEdit PasswordAuthentication no ChallengeResponseAuthentication no PubkeyAuthentication yes`
- Restarted the SSH service

Now, login is only possible using my private key. No key? No access.

Even if someone *did* find my server, there's no username/password prompt to exploit.

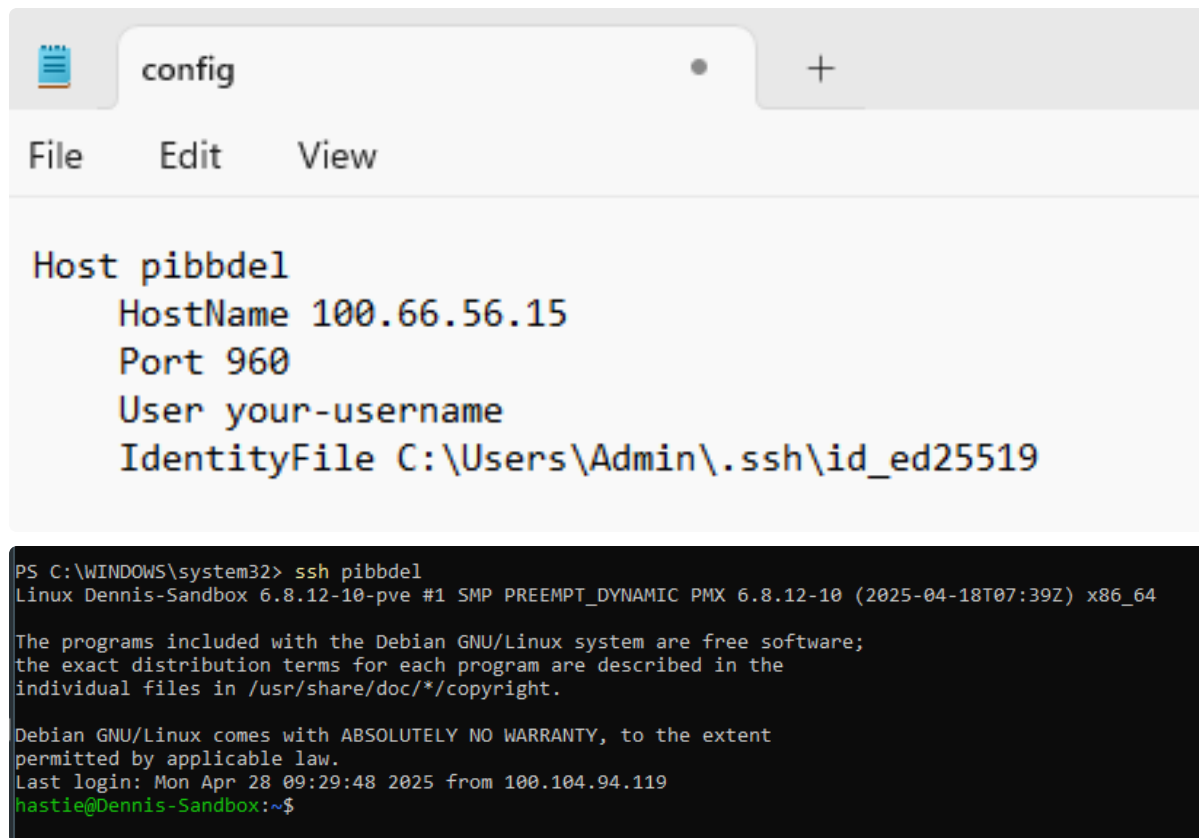
🧠 SSH Config Quality of Life

To avoid typing full SSH commands every time, I created a simple config entry on my local machine:

bash

CopyEdit

```
Host pibbdel HostName 100.64.x.x Port 960 User denney IdentityFile  
~/.ssh/homelab_ed25519
```



The image shows a text editor window titled 'config' with a menu bar (File, Edit, View). The editor contains the following SSH configuration:

```
Host pibbdel  
  HostName 100.66.56.15  
  Port 960  
  User your-username  
  IdentityFile C:\Users\Admin\.ssh\id_ed25519
```

Below the editor is a terminal window showing the execution of the command `ssh pibbdel`. The output shows the user is logged into a Debian GNU/Linux system (Dennis-Sandbox) with the prompt `hastie@Dennis-Sandbox:~$`.

Now I just type `ssh pibbdel` and I'm in.


🔒 Two-Factor Authentication for the Web GUI

SSH wasn't the only thing I wanted secured. Proxmox's web interface is powerful — and dangerous if compromised.

So I enabled **TOTP (Time-Based One-Time Passwords)** for GUI access.

Steps:

1. Datacenter > Realms > pam > Edit → Enable TOTP
2. Datacenter > Two Factor > Add > TOTP
3. Scanned QR with an authenticator app (I use 2FAs Auth)



Now, logging in requires **both** a password and a one-time code from my phone — just like enterprise services.

Enforcing Zero-Trust Access with Tailscale ACLs

Tailscale has a powerful ACL system that lets you define **who can talk to what**. I edited my config so that:

- Only my device (`denneytho@gmail.com`) can reach:
 - Port `8006` (Proxmox Web GUI)
 - Port `960` (SSH)

Access Controls

Define a policy for which devices and users are allowed to connect in your network. [Learn more](#)

```
1 {
2   "acls": [
3     {
4       "action": "accept",
5       "src": ["denneytho@gmail.com"],
6       "dst": ["100.104.94.119:8006"],
7     },
8     {
9       "action": "accept",
10      "src": ["denneytho@gmail.com"],
11      "dst": ["100.113.13.80:960"],
12    },
13    {
14      "action": "accept",
15      "src": ["autogroup:member"],
16      "dst": ["autogroup:self:*"],
17    },
18    {
19      "action": "accept",
20      "src": ["autogroup:member"],
21      "dst": ["tag:exit-node:*"],
22    },
23  ],
24  "ssh": [
25    {
26      "action": "check",
27      "src": ["autogroup:member"],
28      "dst": ["autogroup:self"],
29      "users": ["autogroup:nonroot", "root"],
```

Everyone else — including other Tailscale devices on the network — are denied. No lateral movement, no assumptions.

This is **zero-trust security**, even in a homelab.

Automated VM Backups

Even a well-secured system can fail — hardware crashes, config errors, or fat-fingered deletions.

So I set up scheduled VM backups via Proxmox:

- **Full backups weekly**
- **Incremental backups nightly**
- Stored to local disk (for now)
- Email notifications enabled for failures

Why?

- Quick rollback in case of breakage
 - Safe experimentation without fear of loss
 - Critical in labs that simulate malicious activity or system compromise
-

✅ Final Thoughts

This homelab isn't just a playground — it's a **controlled environment** where I'm simulating enterprise-grade security practices on a personal scale. Everything I did here — VPN tunneling, 2FA, key-based auth, firewalling, zero-trust ACLs — is directly applicable to real-world IT and cybersecurity roles.

It also laid the groundwork for my next project: simulating **Active Directory and Azure hybrid environments** with built-in misconfigurations, insider threats, and privilege escalation paths.

➡️ SOON Up Next: The HollowRoot Lab

Next post: I'll be diving into **Phase 1** of HollowRoot — a red/blue team project where a rogue intern compromises Active Directory and pivots into cloud infrastructure.